

# EFFECTIVE PLATFORM BUILDING WITH KUBERNETES



Wojciech Barczyński - [SMACC.io](#) | [Hypatos.ai](#)  
18 December 2018

# WOJCIECH BARCZYŃSKI

- Lead Software Developer & System Engineer
- Organizer Golang Warsaw Meetup
- Visiting Lecturer at WSB and ALK
- 2019 ➔ Trainings & Consultancy



# STORY

Kubernetes + Go + ...

- **SMACC** - Fintech / ML - since 2017
- **Lyke** - Mobile Fashion app - since 2016

I do not like Infra

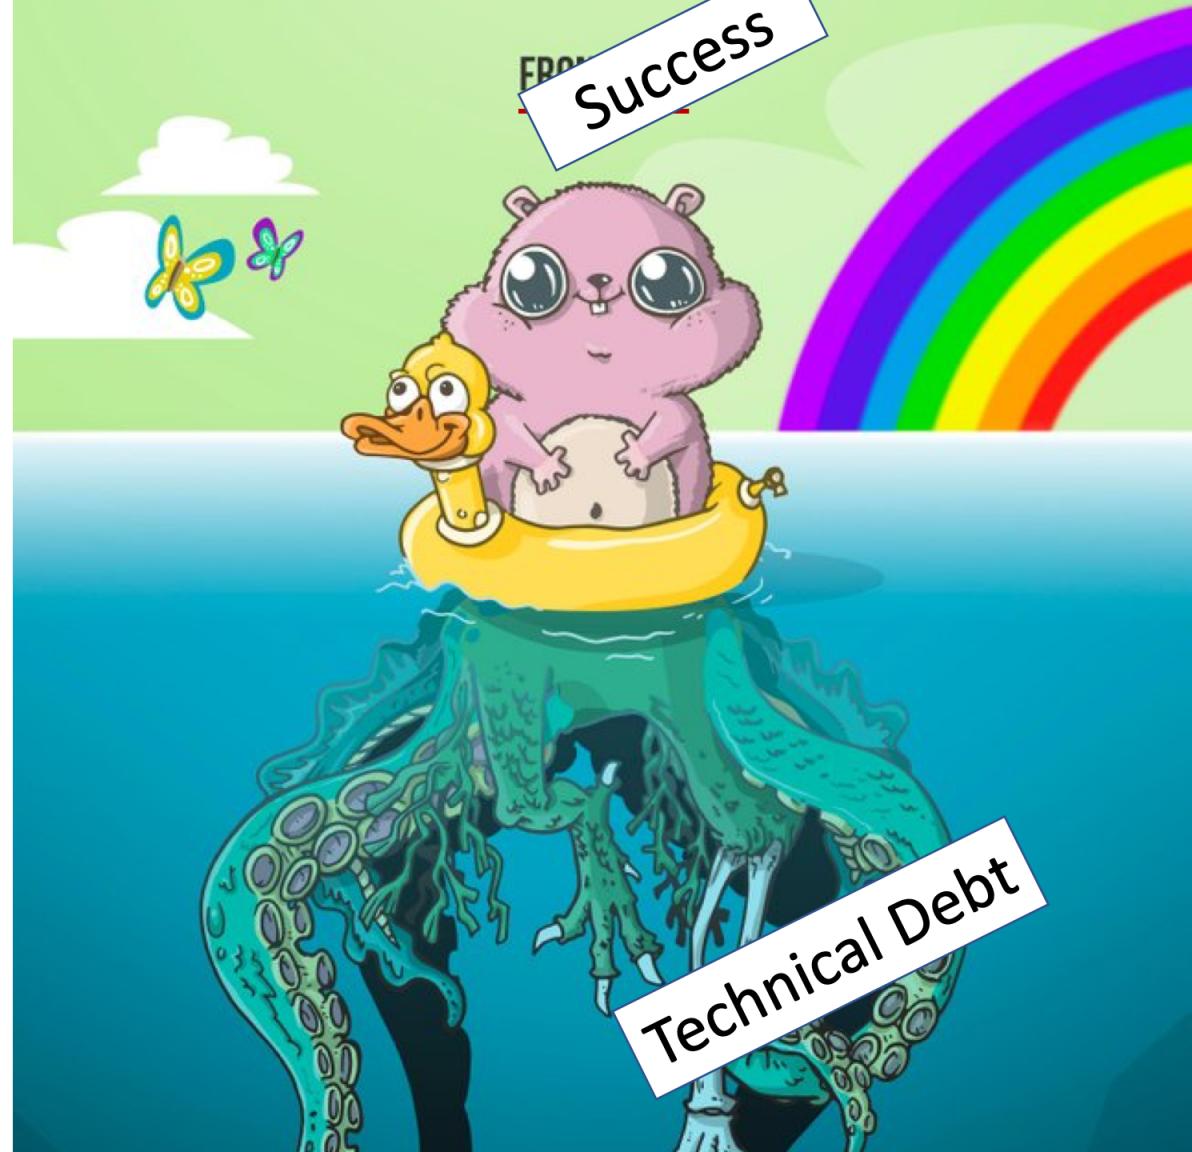
Slow delivery

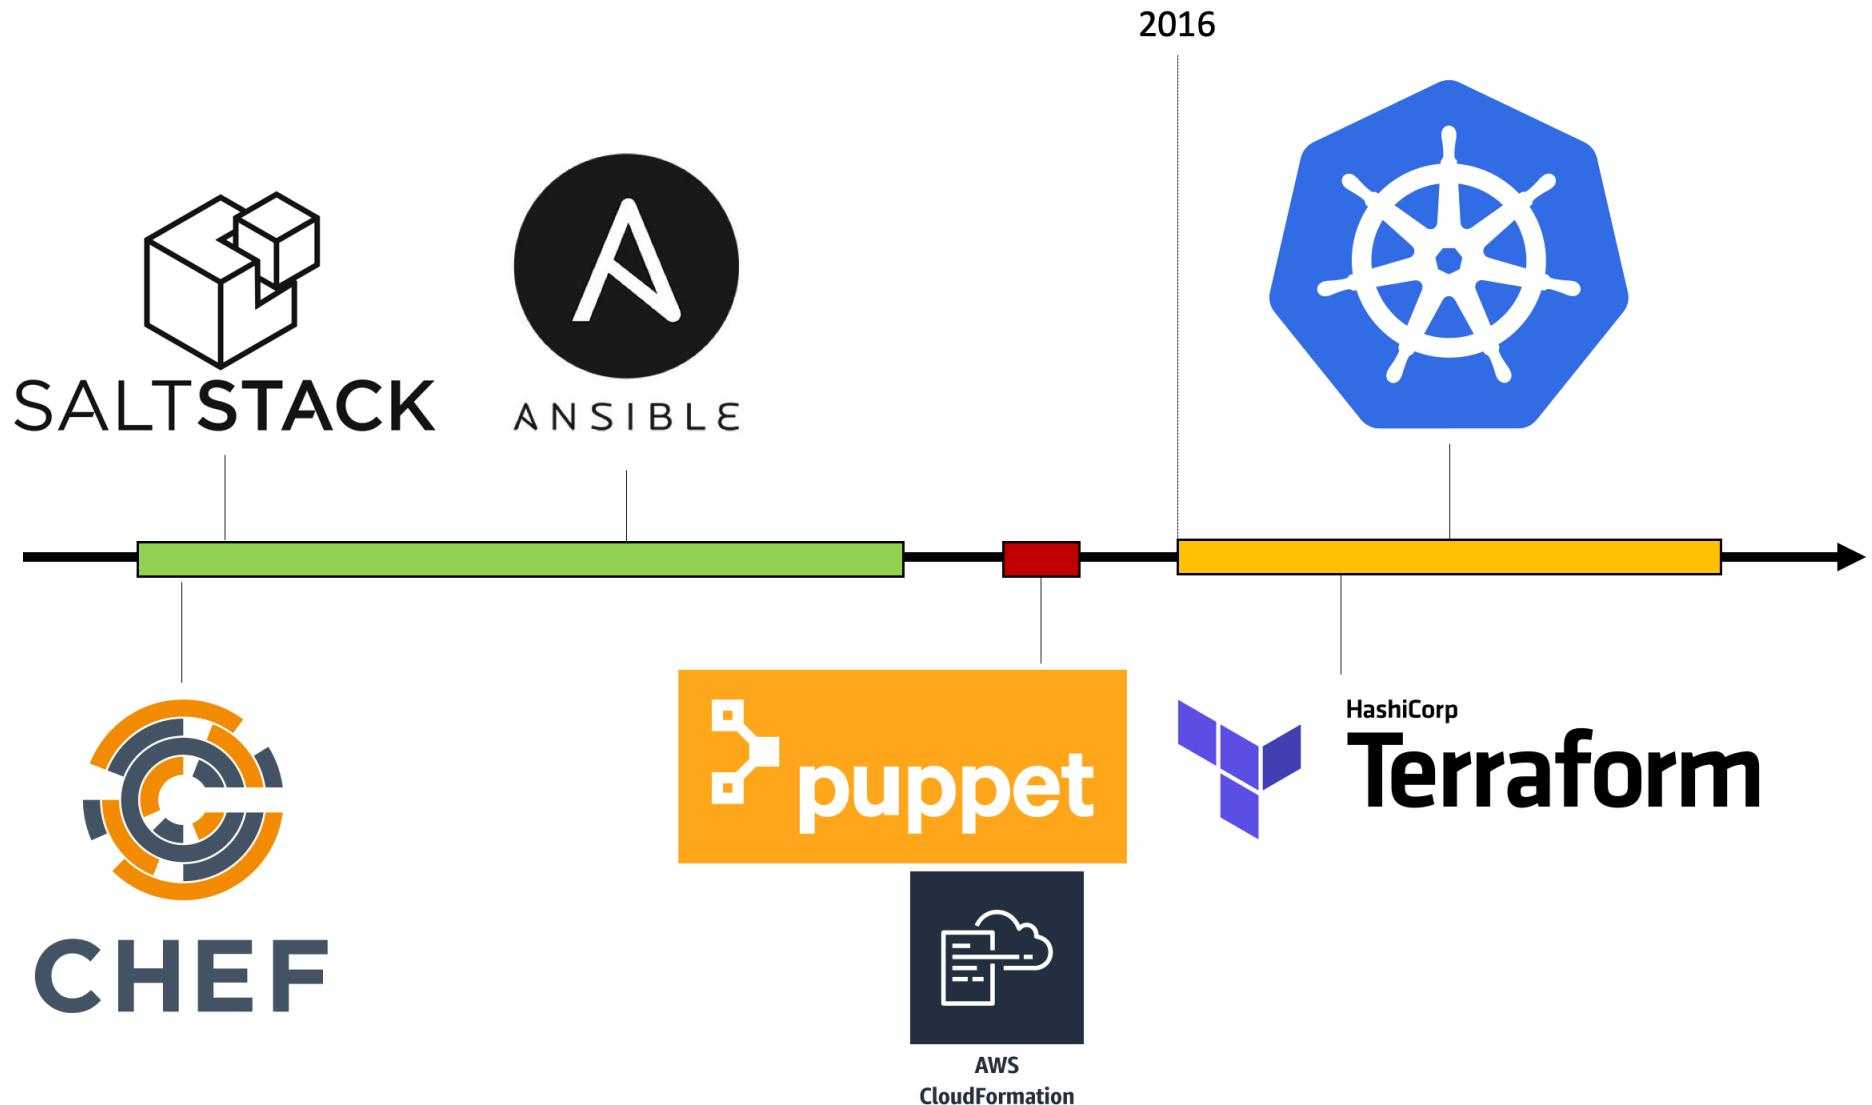
Continuous Deployment?

Fear

Frustration

XX% Idle Machines





# Black (Blue) Box

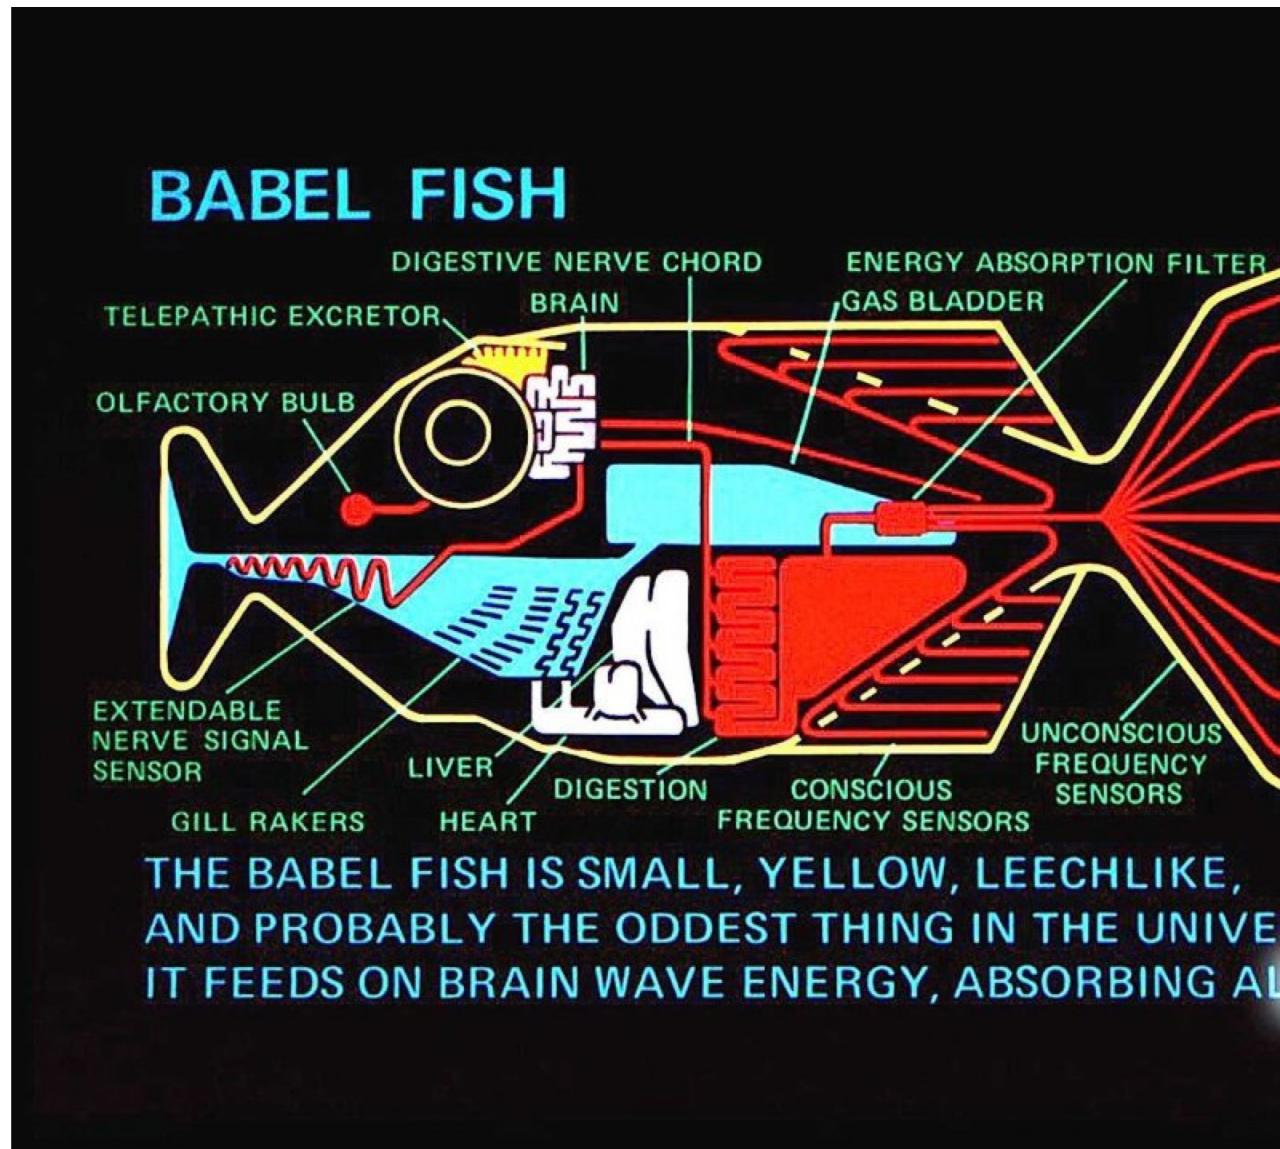
## Infrastructure (almost) invisible

## Easy\* Continuous Deployment



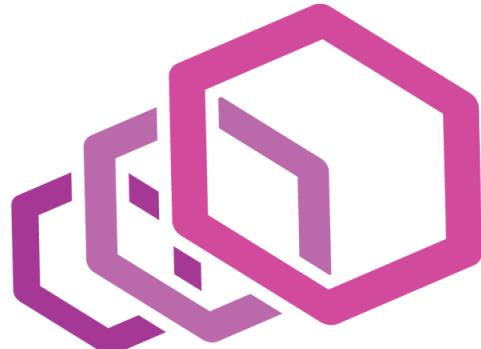
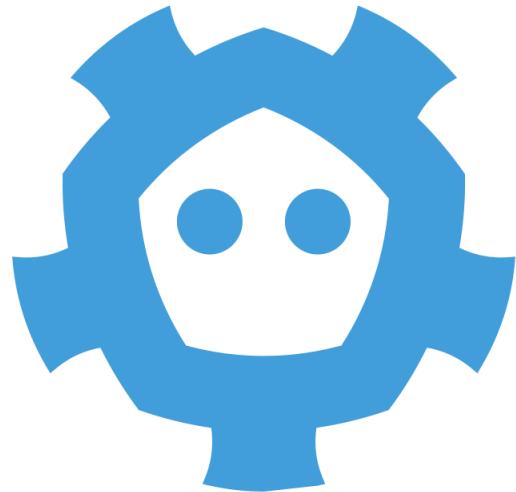
[https://en.wikipedia.org/wiki/File:Dr\\_Who\\_\(316350537\).jpg](https://en.wikipedia.org/wiki/File:Dr_Who_(316350537).jpg)

Common  
Language  
Artifacts  
Platform



# Learn-as-you-go

1. Deploy Cloud-Native app
2. Make a Hell of Mistakes
3. Get it right or Postpone



envoy



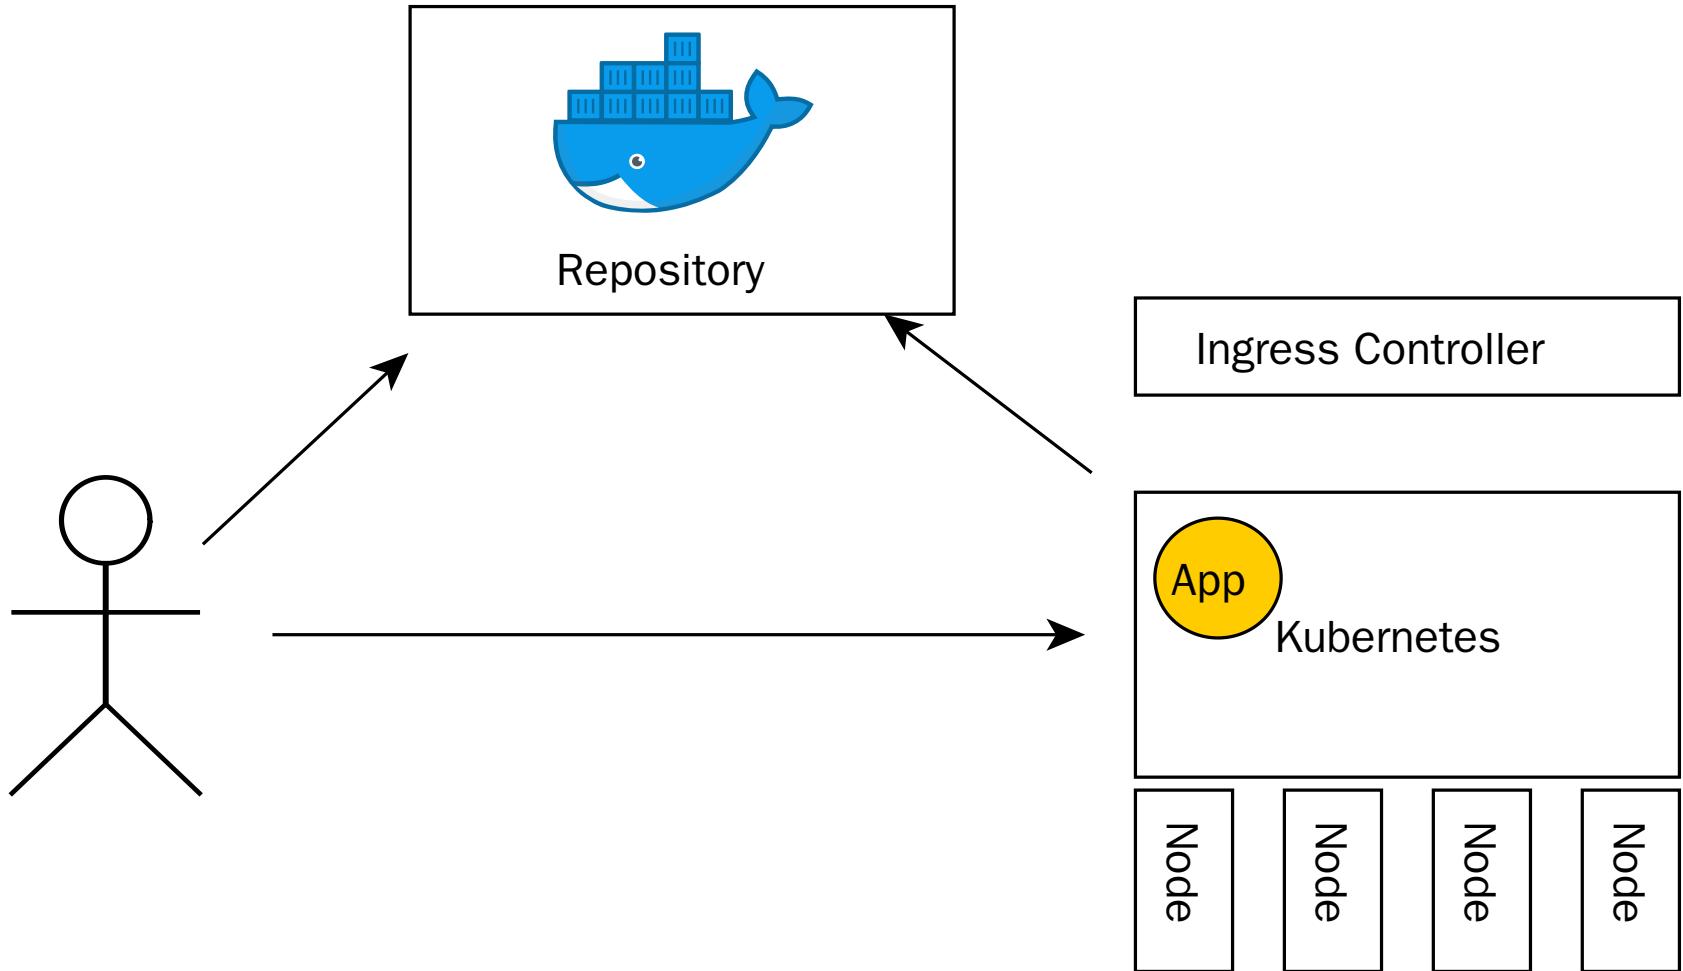
No free lunch - application way smarter  
(12factorapps, coordination, metrics, ...)

<https://www.flickr.com/photos/160866001@N07/>

# KUBERNETES

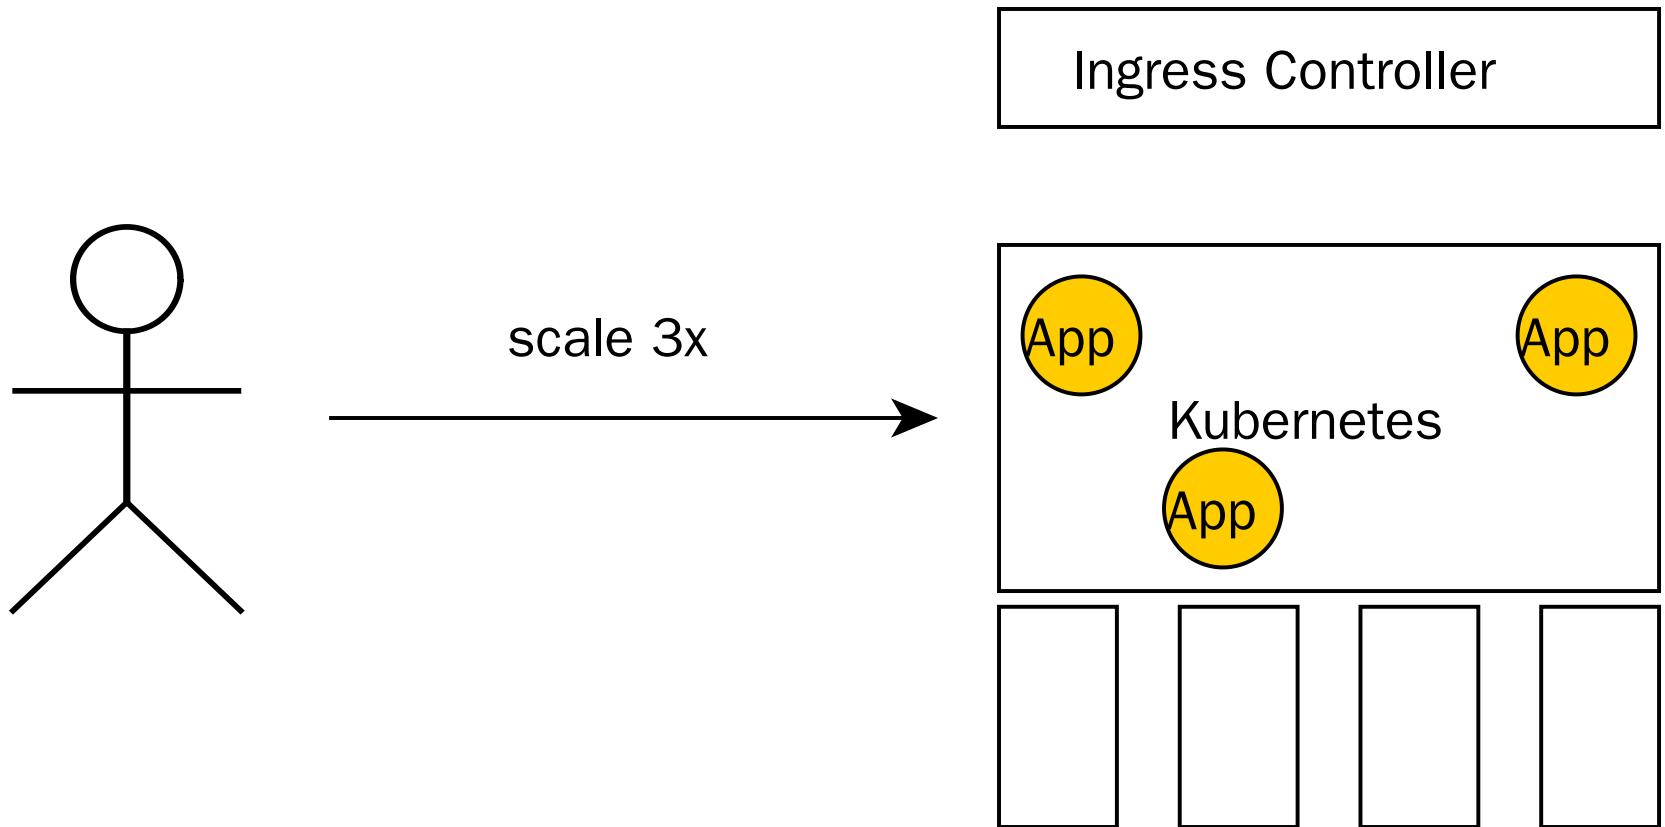
- Container management
- Battery for 12factor apps
- ...heading to integration platform
- ...becoming a framework for your Xubernetes

# KUBERNETES



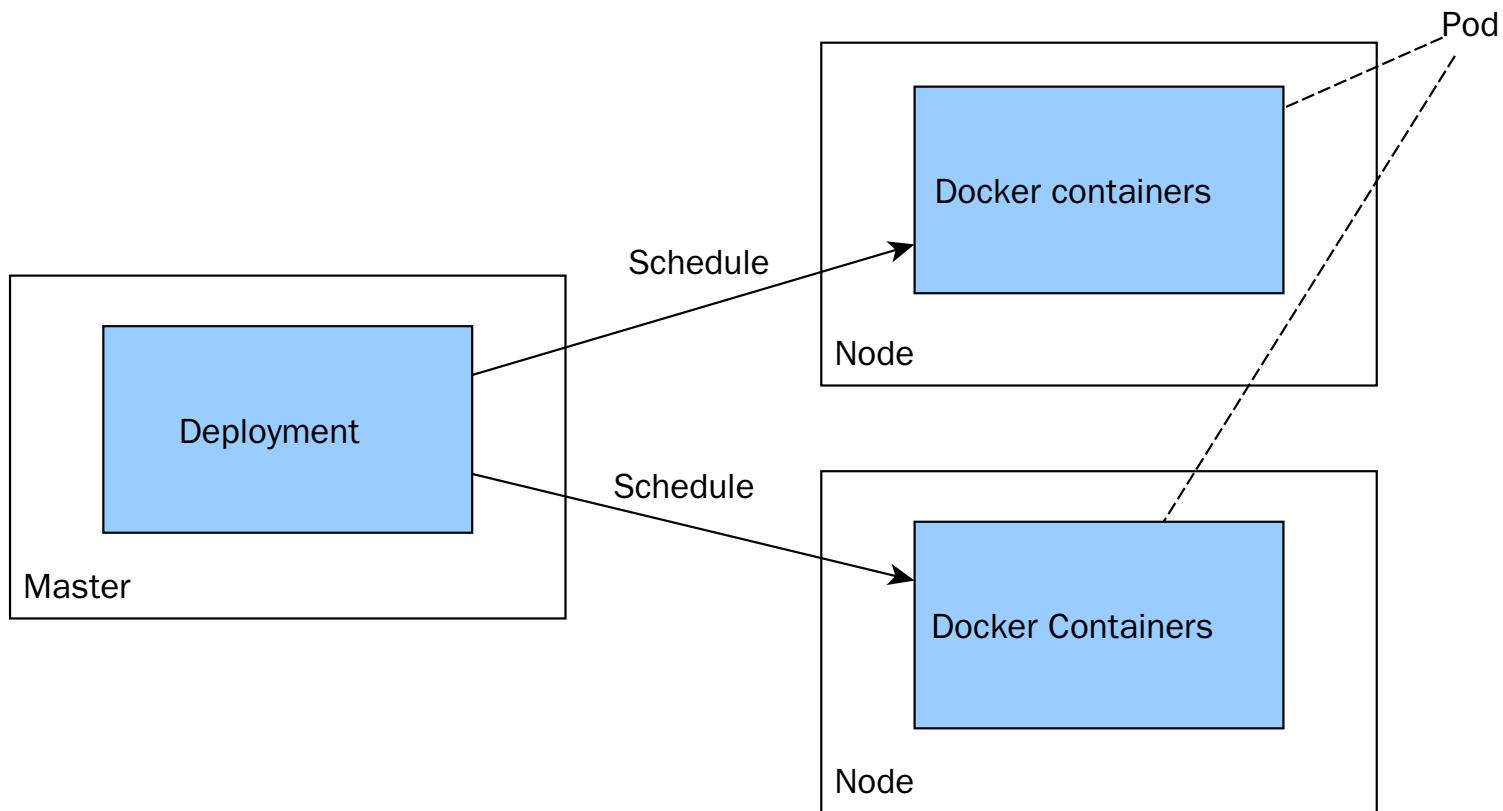
make docker\_push; kubectl create -f app-srv-dpl.yaml

# SCALE UP! SCALE DOWN!

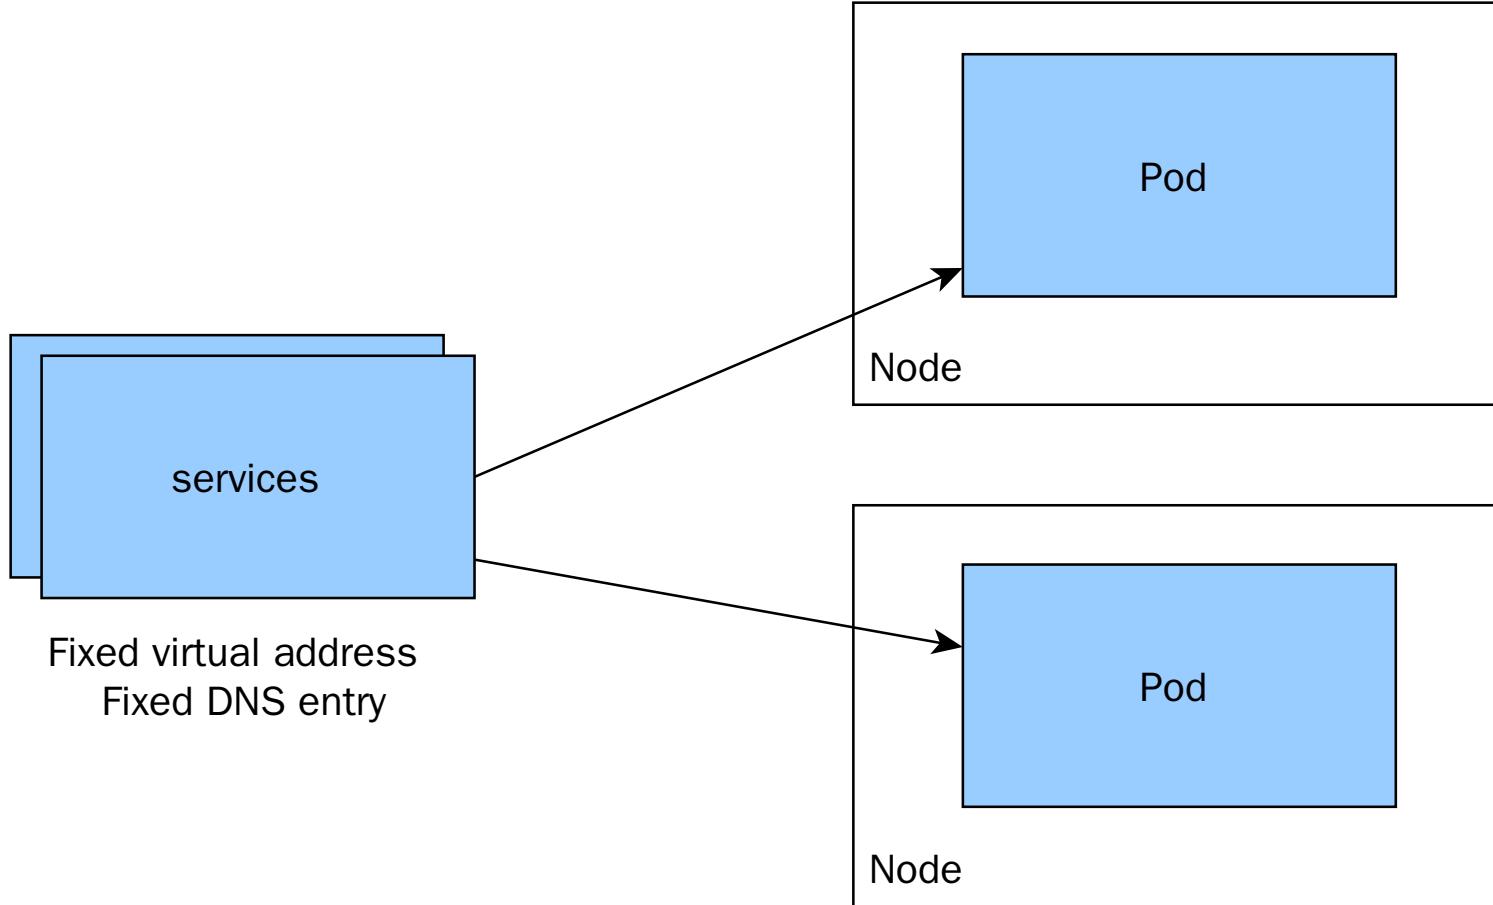


```
kubectl --replicas=3 -f app-srv-dpl.yaml
```

# DEPLOYMENT AND PODS



# SERVICE AND PODS



Service matches pods based on labels

# INGRESS

Pattern

api.smacc.io/v1/users

Target App Service

users-v1

---

api.smacc.io/v2/users

users-v2

---

smacc.io

web

# BASIC CONCEPTS

Name	Purpose	
Service	Interface	Entry point (Service Name)
Deployment	Factory	How many pods, which pods
Pod	Implementation	1+ docker running

# SERVICE DISCOVERY AND LABELING

- names in DNS:

`curl http://users/list`

- labels:

`name=value`

- annotations:

`prometheus.io/scrape: "true"`

# SERVICE DISCOVERY

- loosely couple components
- auto-wiring with logging and monitoring
- drop-in installation (traefik, prometheus, ...)

# ALL HAIL THE GIT!

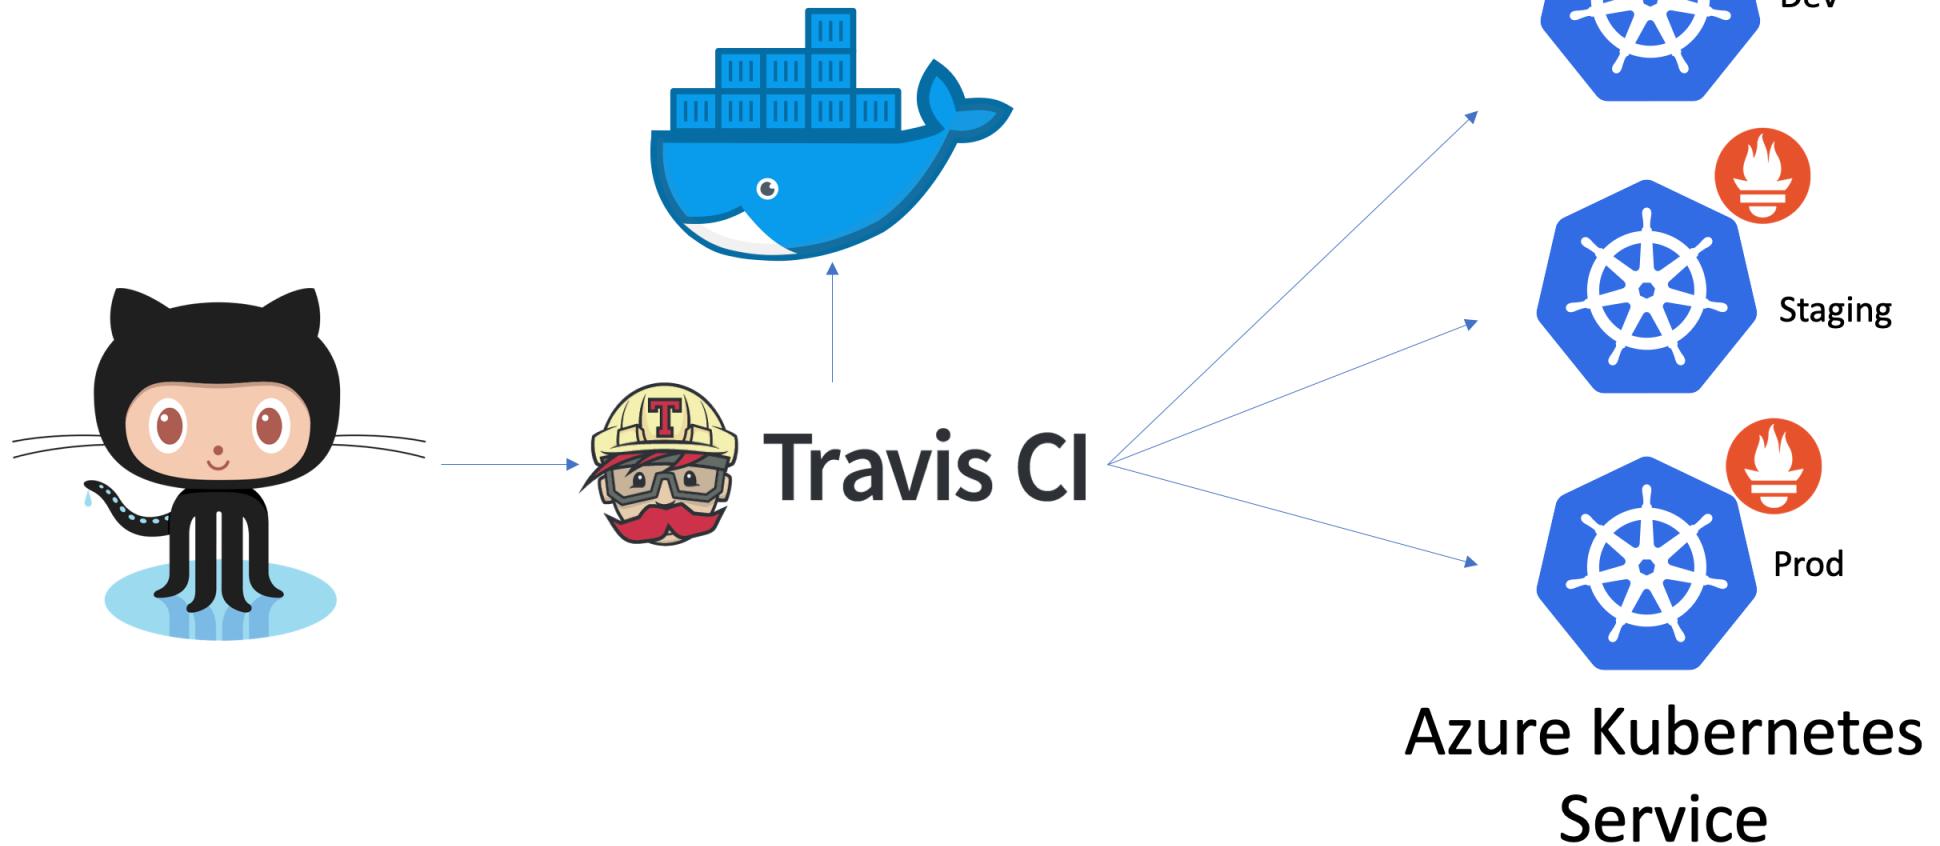
- Yaml
- integration:  
monitoring, alarming, ingress-controller
- ...
- Devs can forget about infra... almost
- DevOps Culture Dream!
  - + all tools -> CRD



Keep it simple

[https://www.flickr.com/photos/bruno\\_brujah/](https://www.flickr.com/photos/bruno_brujah/)

# Continuous Deployment



# Continuous Deployment

1. make run\_local
2. Code on master → develop env



Travis CI



Development

# Continuous Deployment

3. Git tag → staging env
4. PR accepted → production env



Travis CI



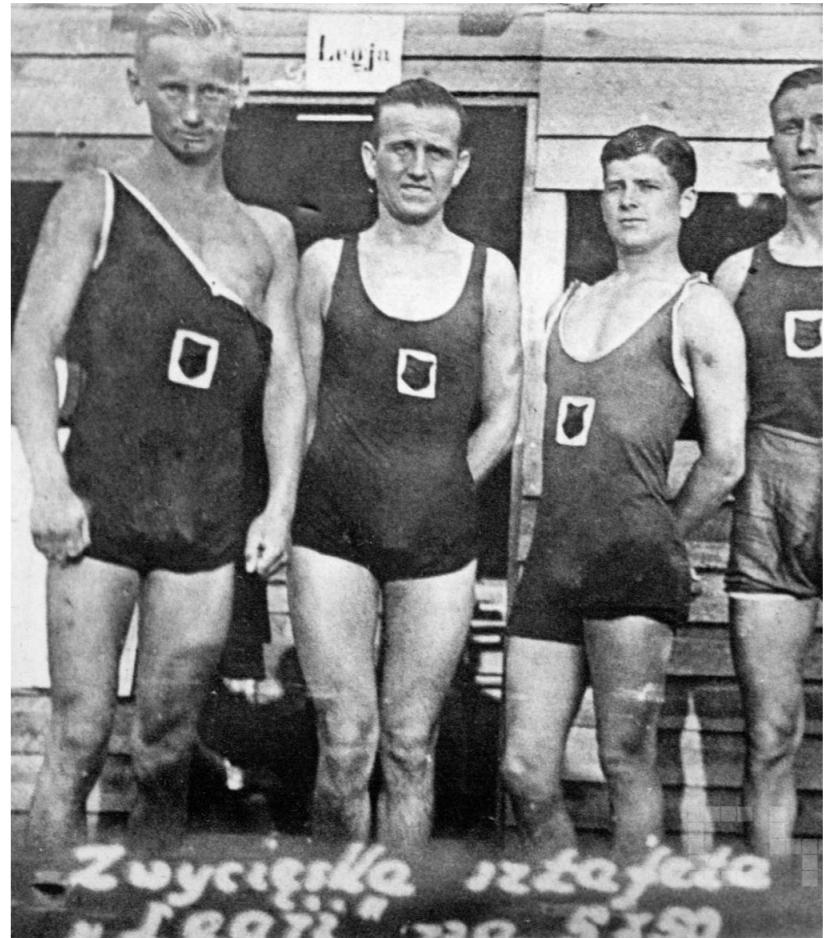
Staging



Production

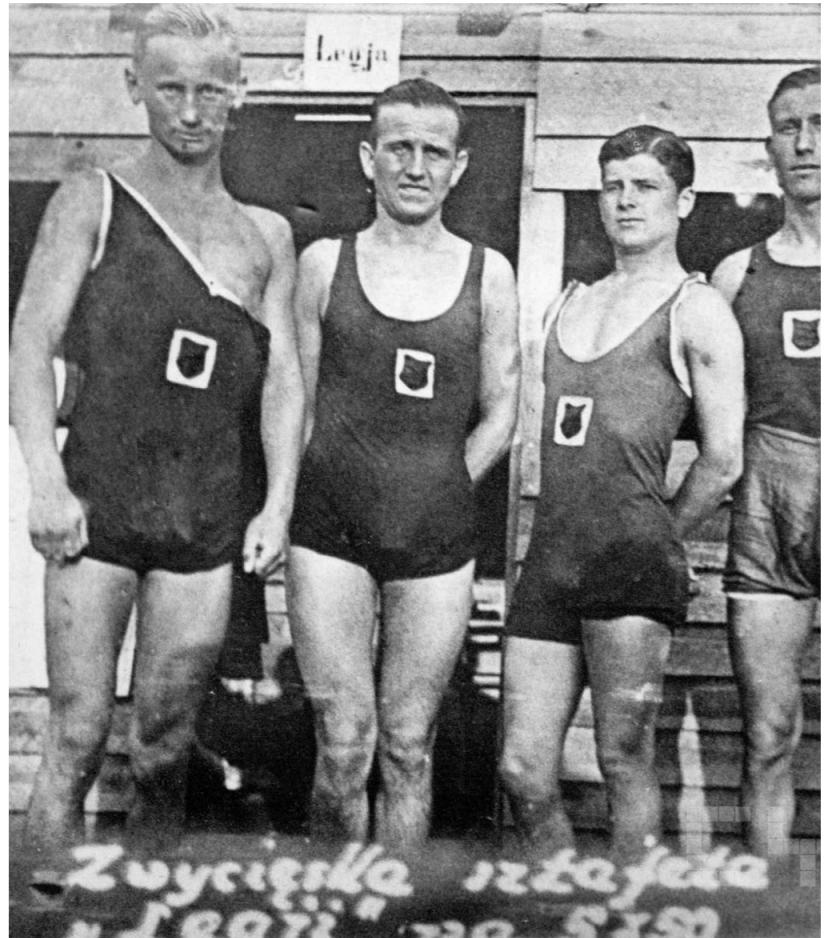
# Keep everybody in the process

1. Teach the team  
Kubernetes definitions
2. Keep the process  
understandable



# Keep everybody in the process

3. Keep K8S files, ...  
easy to read
4. Do not terrorize with  
how-amazing-  
Kubernetes-is :D



# Keep everybody in the process

Copy & Paste:

1. Makefile
2. Kubernetes files
3. TravisCI

```
curl https://github.com/smacc-ci/deploy.sh | bash
```

# Conventions over tools!

- Common conventions for repos
- No a single deploying tool
- No encrypted data in repo

ps. Only when you are really really ready.

# ALL IN GIT

smacc-platform.git

- branches: dev, staging, master
- Generated k8s files
- ENV variables for each of environment

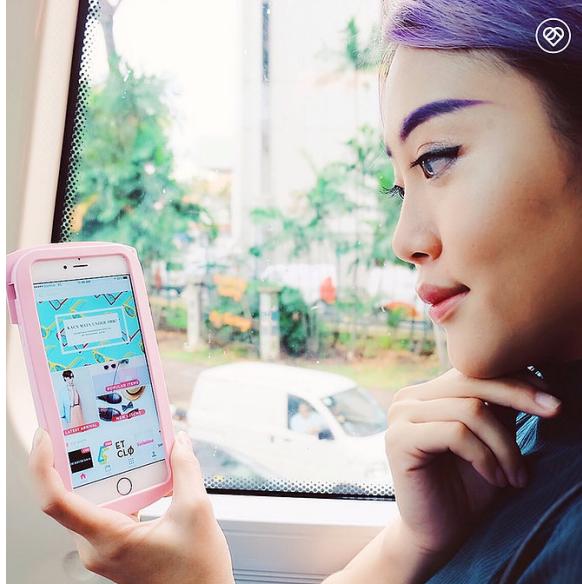
Similar to the [Kelsey Hightower approach](#)

# KUBERNETES

- Pure and generated from simple templates with \${ENV\_VAR}
- 2x kubernetes operators

# LYKE

- E-commerce Indonesia
  - Mobile-only
  - 50k+ users
  - 2M downloads
  - Top 10 Fashion Apps
- Google Play Store



<http://www.news.getlyke.com/single-post/2016/12/02/Introducing-the-New-Beautiful-LYKE>

Now JollyChic Indonesia

## CHALLENGES

- 50+ VMs in Amazon, 1 VM - 1 App
- 65% Idle machine \$\$\$
- Puppet with manual deployment process
- Fear, Forgotten components
- Performance issues

# APPROACH

1. Simplify Infrastructure
2. Change the Development Practices  
(12factor + kubernetes)
3. Make everybody learn Kubernetes
4. Change the Work Organization

see: Conway's law

# SIMPLIFY

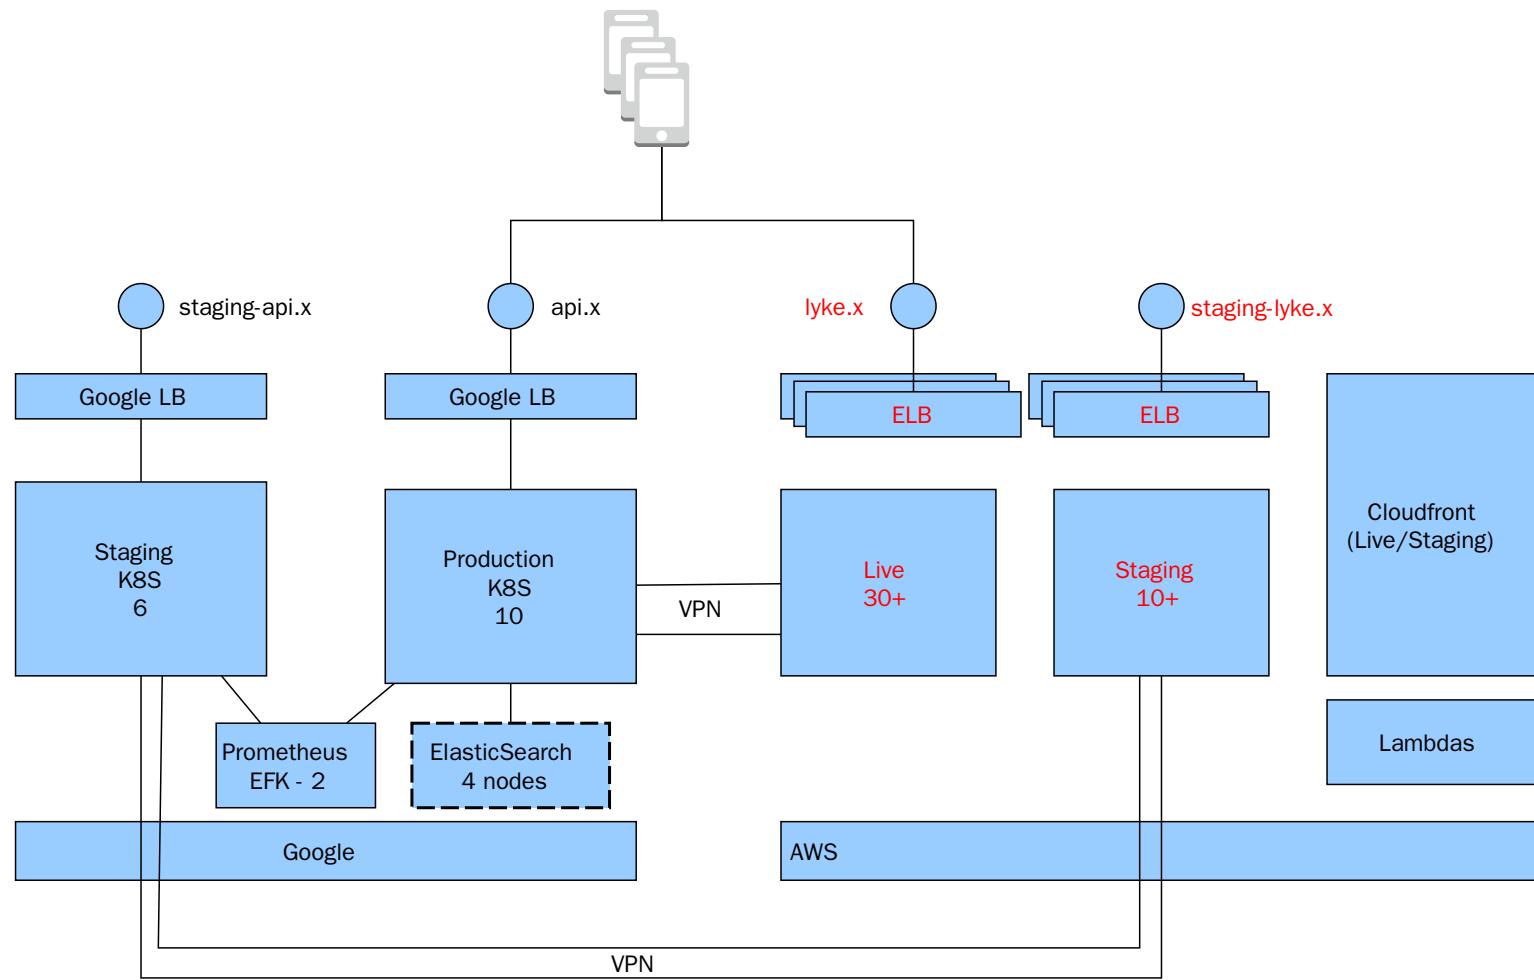
1. Kubernetes with Google Kubernetes Engine
2. Terraform for all new
3. Database-as-a-service

## SIMPLIFY

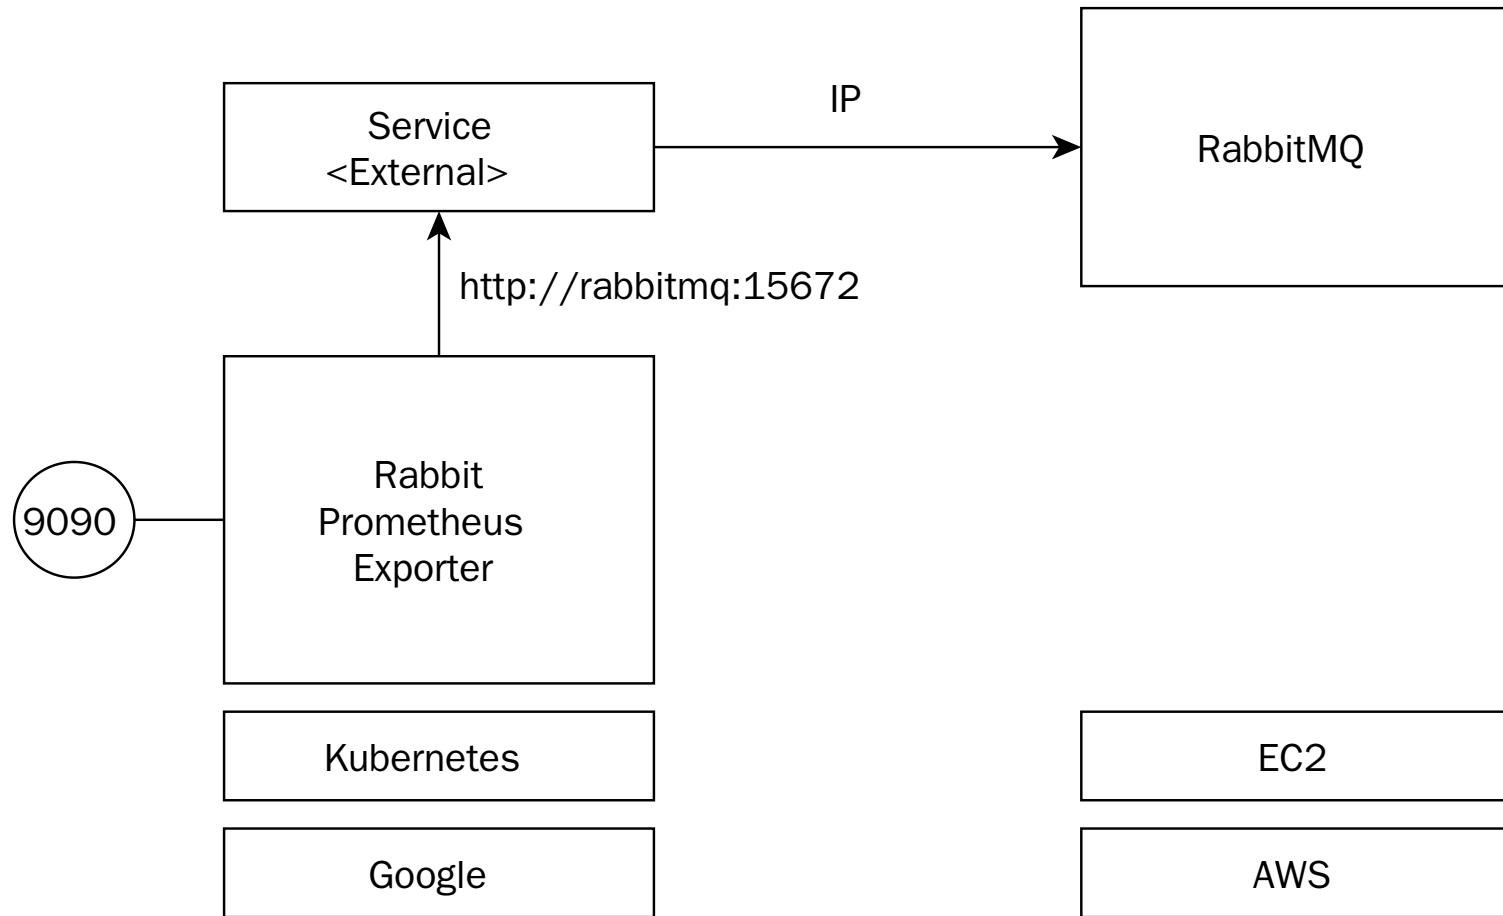
1. Prometheus, AlertManager, and Grafana
2. Elasticsearch-Fluentd-Kibana
3. Google Identity-Aware-Proxy  
to protect all dev dashboards

One person efford

# Architecture During Migration



# Bridge the new with old



Monitor legacy with new stack

# CONTINUOUS DEPLOYMENT

- branch-based:
  - master
  - staging
  - production
- repo independent

Good transition strategy

# CONTINUOUS DEPLOYMENT

- Tests
- Build docker
- Deploy to Google Container Registry and k8s
- No config/secrets applied

Vault planned; moving secrets to runtime

# GIT REPO

```
| - tools
|   | - kube-service.yaml
|   \ - kube-deployment.yaml
|
| - Dockerfile
| - VERSION
\ - Makefile
```

# Makefile

```
SERVICE_NAME=v-connector
GCP_DOCKER_REGISTRY=eu.gcr.io
test: test_short test_integration

run_local:

docker_build: docker_push

kube_create_config:

kube_apply:

kube_deploy:
```

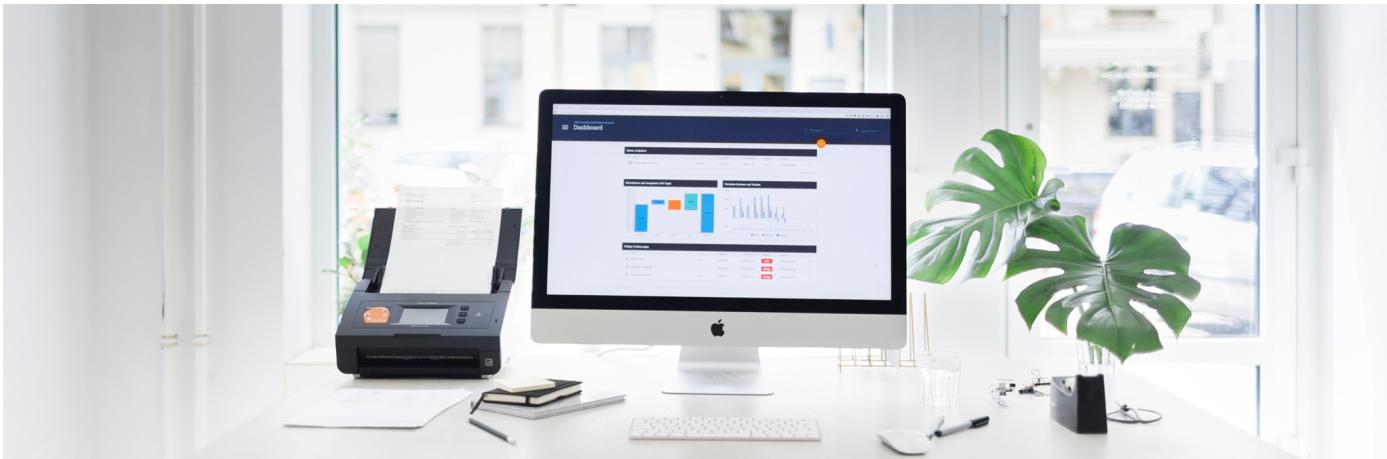
Copy&Paste from the project to project

## WHAT WORKED

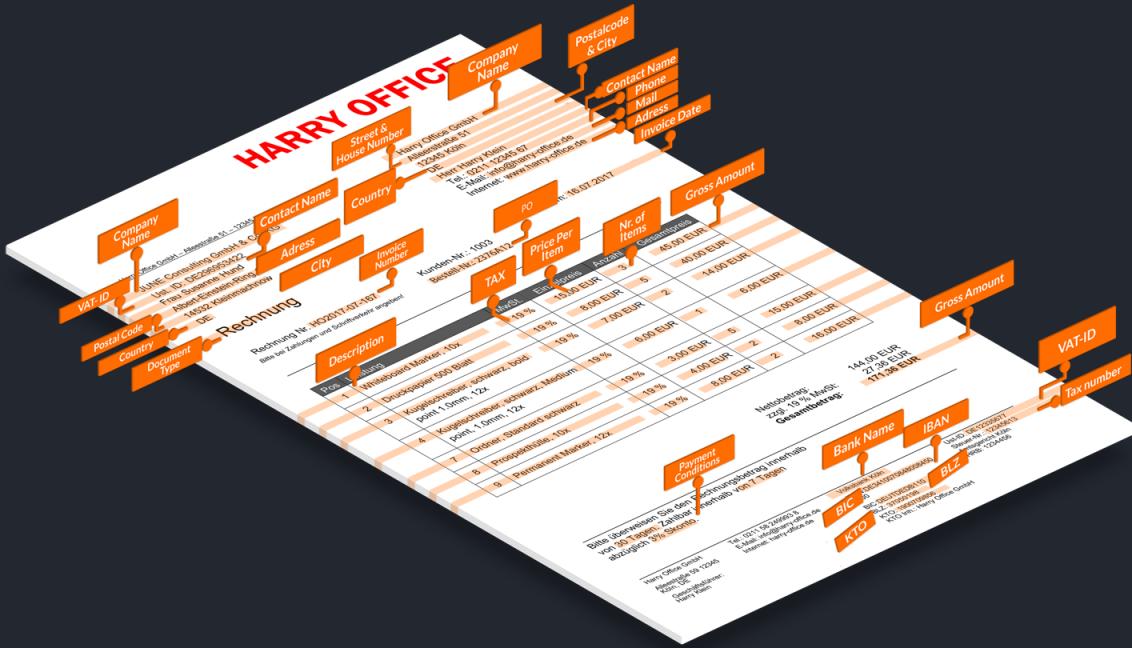
1. Copy&Paste Makefile and k8s yaml
2. Separate deployments a good **transition strategy**

# SMACC

- Machine Learning FinTech
- SaaS and API platform
- From Enterprise (Deutsche Bank, AoK) to SME
- Well-known FinTech Startup in Germany



## Problem SMACC solves



# Hypatos



SMACC

## GET TO K8S

- Legacy on AWS, experiments with AWS ECS :/
- Self-hosted K8S on ProfitBricks
- ooo... K8S can use Azure Active Directory :D

# RUN EVERYWHERE!

- Get to Microsoft ScaleUp, welcome Azure
- Luckily - Azure Kubernetes Service
- With GPU integration :)
- Our OnPrem = Our OnCloud

# KEEP IT SIMPLE

- az aks CLI for setting k8s - README.rst
- Terraform for everything else
- Secrets: 1Password and gopass.pw

Terraform also sets our AWS

# MIGRATION WAS EASY...

- Tech-wise yes
- Team-wise no

# DIFFERENCE



- Two teams in Berlin and Warsaw
- Me in Warsaw

# NEW EXPERIENCE

Team overwhelmed with k8s bits in:

- TravisCI, Makefiles, Yaml's
- feel it is too much hasle

Maybe Keylsey Hightower was right ;)

## APPROACH

- Simplify, Simplify
- Hide K8S magic
- git tag driven Continuous Deployment

# Repo .travis.yml

```
language: go
go:
- '1.16'
services:
- docker
install:
- curl -sL https://$GITHUB_TOKEN@raw.githubusercontent.com
- if [ -f "tools/travis/install.sh" ]; then bash tools/travi
script:
- dep ensure
- make lint
- make test
- if [ -z "${TRAVIS_TAG}" ]; then make snapshot; fi;
deploy:
  provider: docker
  tag: $TRAVIS_TAG
```

# Makefile

```
| - tools
|   | - Makefile
|   | - kube-service.yaml
|   \- kube-deployment.yaml
|
| - Dockerfile
\- Makefile
```

Makefile only tasks for dev

# CONTINUOUS DEPLOYMENT

- Github
- TravisCI
- hub.docker.com
- AKS

# CONTINUOUS DEPLOYMENT

1. git tag and push
2. smacc-platform.git
3. Deploy to staging
4. PR to production

# WHAT WORKED

- Hiding k8s
- Understandable CD process

# **WOULD DO DIFFERENT**

- More sensitive to feedback

# NEXT

- Scale our ML trainings on the top of k8s - kubeflow
- Use CRD to bring exteranl tools to k8s
- Keeping an eye on Istio
- Deployment tool based on [missy](#)

# Effective Kubernetes

- Start with small iterations
- Learn as-you-go
- Keep Kubernetes Understable

Hope the k8s community keep it this way

# Effective Kubernetes

- Move configuration to runtime
- Do not terrorize your devs with K8S
- No free lunch... app must be smarter

# Big thanks to my colleagues in Lyke and SMACC

```
123 def distance_matrix(regions):  
124     """ Computes a distance matrix against a region list """  
125     tuples = [r.as_tuple() for r in regions]  
126     return cdist(tuples, tuples, region_distance)  
127  
128  
129 def clusterize(words, **kwargs):  
130     # TODO: write a cool docstring here  
131     db = DBSCAN(metric="precomputed", **kwargs)  
132     X = distance_matrix([Region.from_word(w) for w in words])  
133     labels = [int(l) for l in db.fit_predict(X)]
```



# THANK YOU. QUESTIONS?

```
123 def distance_matrix(regions):  
124     """ Computes a distance matrix against a region list """  
125     tuples = [r.as_tuple() for r in regions]  
126     return cdist(tuples, tuples, region_distance)  
127  
128  
129 def clusterize(words, **kwargs):  
130     # TODO: write a cool docstring here  
131     db = DBSCAN(metric="precomputed", **kwargs)  
132     X = distance_matrix([Region.from_word(w) for w in words])  
133     labels = [int(l) for l in db.fit_predict(X)]
```



[github.com/wojciech12/talk\\_effective\\_kubernetes](https://github.com/wojciech12/talk_effective_kubernetes)

```
123 def distance_matrix(regions):  
124     """ Computes a distance matrix against a region list """  
125     tuples = [r.as_tuple() for r in regions]  
126     return cdist(tuples, tuples, region_distance)  
127  
128  
129 def clusterize(words, **kwargs):  
130     # TODO: write a cool docstring here  
131     db = DBSCAN(metric="precomputed", **kwargs)  
132     X = distance_matrix([Region.from_word(w) for w in words])  
133     labels = [int(l) for l in db.fit_predict(X)]
```



# SMACC



Go



PYTORCH

TensorFlow™



amazon  
web services™



Azure



# BACKUP SLIDES

```
123 def distance_matrix(regions):  
124     """ Computes a distance matrix against a region list """  
125     tuples = [r.as_tuple() for r in regions]  
126     return cdist(tuples, tuples, region_distance)  
127  
128  
129 def clusterize(words, **kwargs):  
130     # TODO: write a cool docstring here  
131     db = DBSCAN(metric="precomputed", **kwargs)  
132     X = distance_matrix([Region.from_word(w) for w in words])  
133     labels = [int(l) for l in db.fit_predict(X)]
```



# INFRA TOOLS

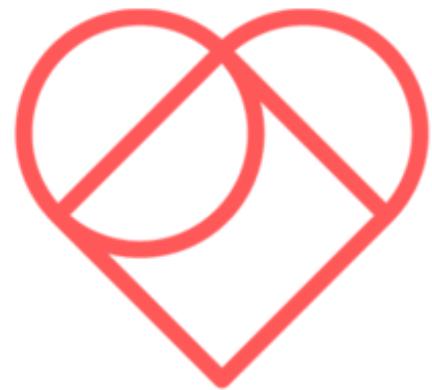
- Prometheus + AlertMnager + Grafana
- Traefik
- Kafka - Yolean/kubernetes-kafka
- Vault on Etcd - banzaicloud/bank-vaults

# DATA STORES

- Kafka - Yolean/kubernetes-kafka
- Etcd - coreos/etcd-operator
- DB: PSQL and Mongo

# **BACKUPS:**

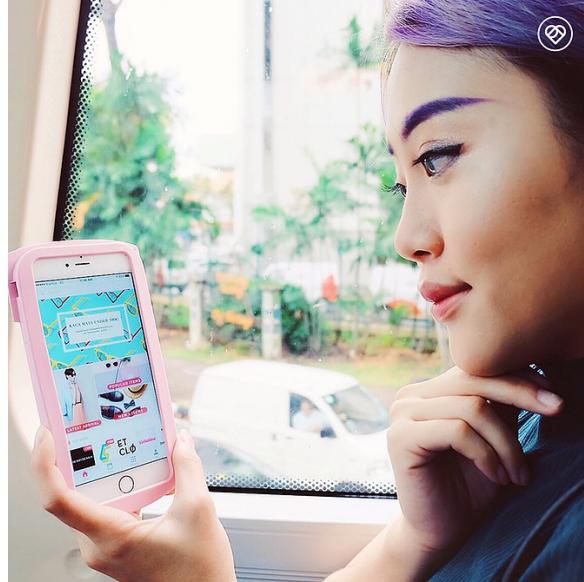
- old-school backups with ARK and Restic
- replications across clouds



LYKE

# LYKE

- E-commerce
- Mobile-only
- 50k+ users
- 2M downloads
- Top 10 Fashion Apps  
w Google Play Store



<http://www.news.getlyke.com/single-post/2016/12/02/Introducing-the-New-Beautiful-LYKE>

Now JollyChic Indonesia

# GOOD PARTS

- Fast Growth
- A/B Testing
- Data-driven
- Product Manager,  
UI Designer,  
Mobile Dev,  
and tester - one  
body



# 1. CLEAN UP

- Single script for repo - Makefile [1]
- Resurrect the README

[1] With zsh or bash auto-completion plug-in in your terminal.

## 2. GET BACK ALL THE KNOWLEDGE

Extract from:

- Puppet, ... ➔ Dockerfile
- Running Instances ➔ Dockerfile, README.rst
- Nagios, ... ➔ README.rst + *checks/*

## 3. INTRODUCE RUN\_LOCAL

- make run\_local
- A nice section on how to run in README.rst
- with docker-compose

The most crucial point.

## 4. GET TO KUBERNETES

- make kube\_create\_config
- make kube\_apply
- Generate the yaml files if your envs differ  
secrets from gopass (password manager)

## 5. CONTINUOUS DEPLOYMENT

Travis:

- the same Makefile as dev use

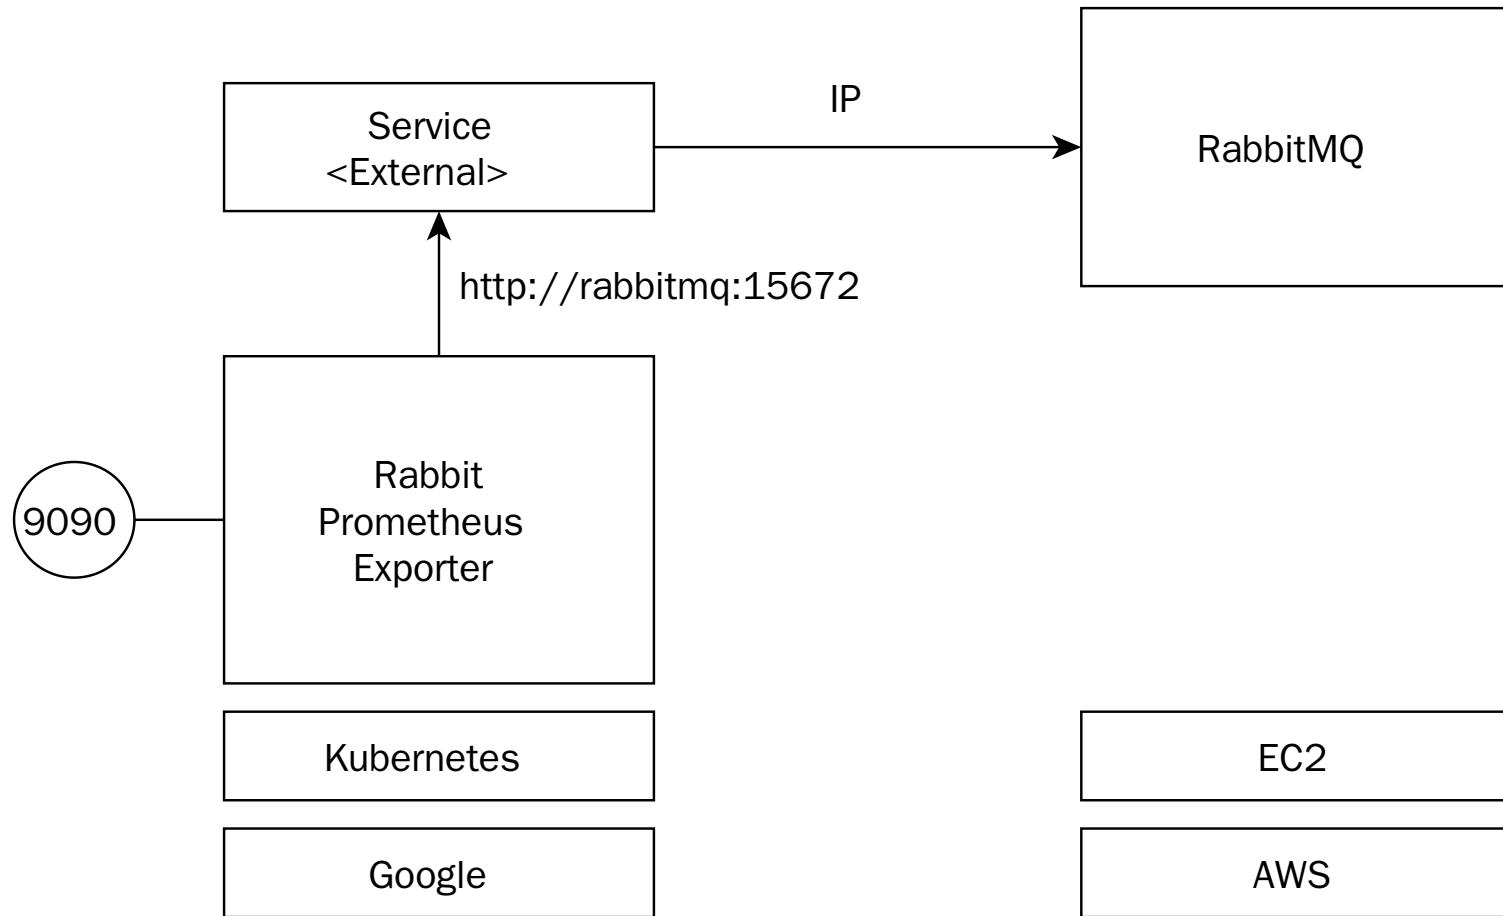
## 6. KEEP IT RUNNING

Bridge the new with old:

- Use External Services in Kubernetes
- Expose k8s in the Legacy [1]

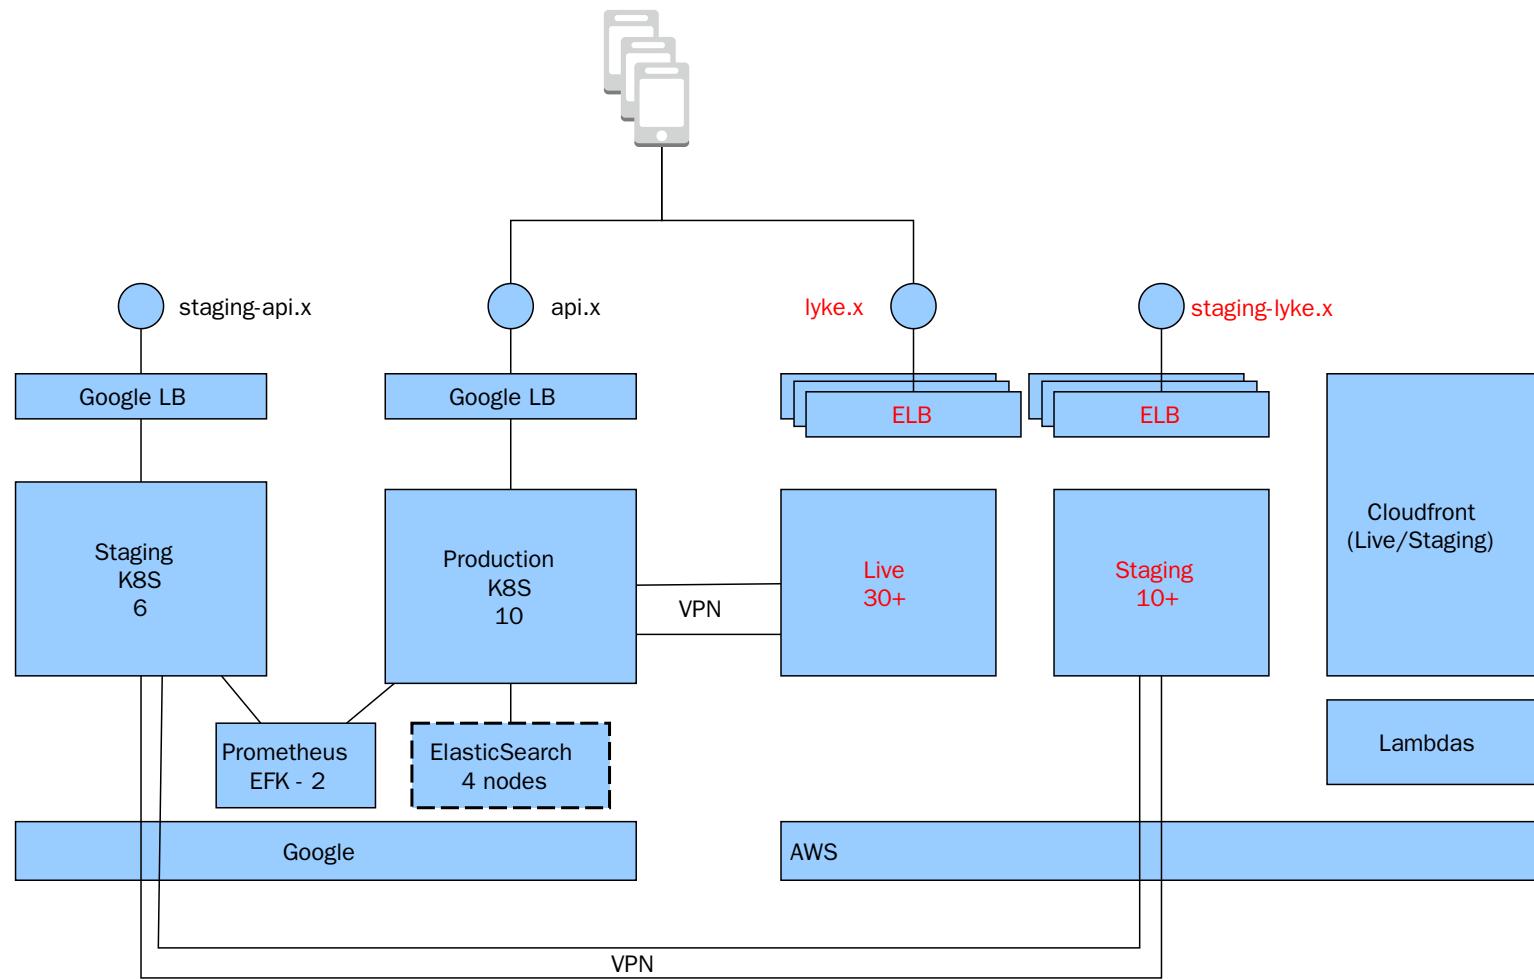
[1] hard coded IP:PORT, considered: K8S events to Consul

# Bridge the new with old



Monitor legacy with new stack

# Architecture During Migration



## 7. INTRODUCE SMOKE-TEST

```
TARGET_URL=127.0.0 make smoke_test
```

```
TARGET_URL=api.example.com/users make smoke_test
```

## 8. SERVICE SELF-AWARE

Add to old services:

1. *metrics/*
2. *health/*
3. *info/*

## 9. MOVE TO MICRO-SERVICES

Offload Legacy:

- Keep the lights on
- New functionality to micro-services

## WHAT WORKED

1. Copy&Paste Makefile and k8s yaml
2. Separate deployments a good transition strategy

# WHAT DID NOT WORK

1. Too many PoC, cut to 2 weeks max
2. Doing it with too small steps
3. Push back to k8s yaml [\*]
4. Alert rules too hard to write

[\*] With coaching, I thought, it is OK

# DO DIFFERENT

1. Move data day one
2. Make devs know it is a transition stage
3. Teach earlier about resources
4. EFK could wait
5. All-hands for a paid-XXX% weekend for migration

# SMACC

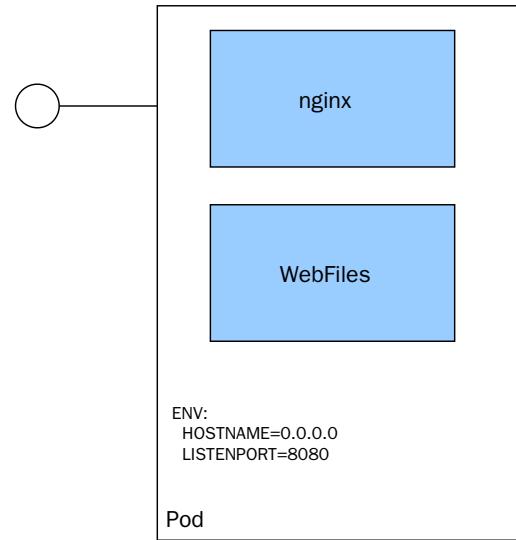
---

## Hypatos

# KUBERNETES CONCEPTS+

# PODS

- See each other on localhost
- Live and die together
- Can expose multiple ports



# SIDE-CARS

